

# assignment3

October 20, 2023

## 1 Assignment 3 - Building a Custom Visualization

In this assignment you must choose **one** of the options presented below and submit a visual as well as your source code for peer grading. The details of how you solve the assignment are up to you, although your assignment must use matplotlib so that your peers can evaluate your work. The options differ in challenge level, but there are no grades associated with the challenge level you chose. However, your peers will be asked to ensure you at least met a minimum quality for a given technique in order to pass. Implement the technique fully (or exceed it!) and you should be able to earn full grades for the assignment.

Ferreira, N., Fisher, D., & Konig, A. C. (2014, April). [Sample-oriented task-driven visualizations: allowing users to make better, more confident decisions](#). In Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (pp. 571-580). ACM. ([video](#))

In this [paper](#) the authors describe the challenges users face when trying to make judgements about probabilistic data generated through samples. As an example, they look at a bar chart of four years of data (replicated below in Figure 1). Each year has a y-axis value, which is derived from a sample of a larger dataset. For instance, the first value might be the number votes in a given district or riding for 1992, with the average being around 33,000. On top of this is plotted the 95% confidence interval for the mean (see the boxplot lectures for more information, and the `yerr` parameter of `barcharts`).

Figure 2c from (Ferreira et al. 2014). Note that the colorbar legend at the bottom as well as the arrows are not required in the assignment descriptions below.

A challenge that users face is that, for a given y-axis value (e.g. 42,000), it is difficult to know which x-axis values are most likely to be representative, because the confidence levels overlap and their distributions are different (the lengths of the confidence interval bars are unequal). One of the solutions the authors propose for this problem (Figure 2c) is to allow users to indicate the y-axis value of interest (e.g. 42,000) and then draw a horizontal line and color bars based on this value. So bars might be colored red if they are definitely above this value (given the confidence interval), blue if they are definitely below this value, or white if they contain this value.

Figure 2c from (Ferreira et al. 2014). Note that the colorbar legend at the bottom as well as the arrows are not required in the assignment descriptions below.

**Easiest option:** Implement the bar coloring as described above - a color scale with at least three colors, (e.g. blue, white, and red). Assume the user provides the y axis value of interest as a parameter or variable.

**Harder option:** Implement the bar coloring as described in the paper, where the color of the bar

is actually based on the amount of data covered (e.g. a gradient ranging from dark blue for the distribution being certainly below this y-axis, to white if the value is certainly contained, to dark red if the value is certainly not contained as the distribution is above the axis).

**Even Harder option:** Add interactivity to the above, which allows the user to click on the y axis to set the value of interest. The bar colors should change with respect to what value the user has selected.

**Hardest option:** Allow the user to interactively set a range of y values they are interested in, and recolor based on this (e.g. a y-axis band, see the paper for more details).

---

*Note: The data given for this assignment is not the same as the data used in the article and as a result the visualizations may look a little different.*

```
[1]: # Use the following data for this assignment:
```

```
import pandas as pd
import numpy as np

np.random.seed(12345)

df = pd.DataFrame([np.random.normal(32000,200000,3650),
                    np.random.normal(43000,100000,3650),
                    np.random.normal(43500,140000,3650),
                    np.random.normal(48000,70000,3650)],
                    index=[1992,1993,1994,1995])

df
```

```
[1]:
```

	0	1	2	3	\
1992	-8941.531897	127788.667612	-71887.743011	-79146.060869	
1993	-51896.094813	198350.518755	-123518.252821	-129916.759685	
1994	152336.932066	192947.128056	389950.263156	-93006.152024	
1995	-69708.439062	-13289.977022	-30178.390991	55052.181256	

	4	5	6	7	\
1992	425156.114501	310681.166595	50581.575349	88349.230566	
1993	216119.147314	49845.883728	149135.648505	62807.672113	
1994	100818.575896	5529.230706	-32989.370488	223942.967178	
1995	152883.621657	12930.835194	63700.461932	64148.489835	

	8	9	...	3640	3641	\
1992	185804.513522	281286.947277	...	171938.760289	150650.759924	
1993	23365.577348	-109686.264981	...	-44566.520071	101032.122475	
1994	-66721.580898	47826.269111	...	165085.806360	74735.174090	
1995	-29316.268556	59645.677367	...	-13901.388118	50173.686673	

	3642	3643	3644	3645	\
1992					
1993					
1994					
1995					

1992	203663.976475	-377877.158072	-197214.093861	24185.008589
1993	117648.199945	160475.622607	-13759.888342	-37333.493572
1994	107329.726875	199250.734156	-36792.202754	-71861.846997
1995	53965.990717	4128.990173	72202.595138	39937.199964

	3646	3647	3648	3649
1992	-56826.729535	-67319.766489	113377.299342	-4494.878538
1993	103019.841174	179746.127403	13455.493990	34442.898855
1994	26375.113219	-29328.078384	65858.761714	-91542.001049
1995	139472.114293	59386.186379	73362.229590	28705.082908

[4 rows x 3650 columns]

```
[2]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 4 entries, 1992 to 1995
Columns: 3650 entries, 0 to 3649
dtypes: float64(3650)
memory usage: 114.1 KB
```

```
[3]: df[df.index == 1992]
```

```
[3]:
```

	0	1	2	3	4	\
1992	-8941.531897	127788.667612	-71887.743011	-79146.060869	425156.114501	
	5	6	7	8	9	\
1992	310681.166595	50581.575349	88349.230566	185804.513522	281286.947277	
...	3640	3641	3642	3643	\	
1992	...	171938.760289	150650.759924	203663.976475	-377877.158072	
	3644	3645	3646	3647	3648	\
1992	-197214.093861	24185.008589	-56826.729535	-67319.766489	113377.299342	
	3649					
1992	-4494.878538					

[1 rows x 3650 columns]

```
[4]: %matplotlib inline
import matplotlib.pyplot as plt
```

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

np.random.seed(12345)

data = [
    np.random.normal(32000, 200000, 3650),
    np.random.normal(43000, 100000, 3650),
    np.random.normal(43500, 140000, 3650),
    np.random.normal(48000, 70000, 3650)
]

years = [1992, 1993, 1994, 1995]

df = pd.DataFrame(data, index=years)

means = df.mean(axis=1)
std_errors = df.sem(axis=1) * 1.96

fig, ax = plt.subplots()

bars = ax.bar(years, height = means, yerr=std_errors, align='center', alpha=0.
    ↪7, capsize=10, color = ['navy', 'bisque', 'dodgerblue', 'maroon'], edgecolor =
    ↪'black',
    width =1)

ax.set_xlabel('Year')
ax.set_ylabel('MeanValue')
ax.set_title('Mean Values of Sample Dataset with Error Bars')
ax.set_yticks(range(0,55000,6500))


reference_y = 42000
ax.axhline(reference_y, color='gray', linestyle='--')

norm = plt.Normalize(means.min(), means.max())
sm = plt.cm.ScalarMappable(cmap="coolwarm", norm=norm)
sm.set_array([])
cbar = plt.colorbar(sm, ax=ax)

ax.set_xticks(years)

```

```
plt.show()
```



```
[ ]:
```